

## Bayesian Statistics

### Exercises 4.

1. Consider the following setting in the context of “multiple hypothesis tests”. Let  $i = 1, \dots, n$  index individuals and  $j = 1, \dots, m$  index genes (or pixels in an image if you prefer). Assume we have measurements on each individual at each gene for “treatments”  $k = 0, 1$ . Let  $Y_{ijk}$  denote the measurement on individual  $i$ , gene  $j$ , treatment  $k$ , and let  $D_{ij}$  denote the difference between the measurements in the two treatments:  $D_{ij} := Y_{ij0} - Y_{ij1}$ . We will assume that  $D$  is sufficient for all our inferences, and so you can forget about  $Y$  now and work only with  $D$ : I just wanted you to understand where  $D$  might come from in principle. We will assume a model for  $D$ :

$$D_{ij} | \beta, \sigma \sim N(\beta_j, \sigma_j^2) \quad (1)$$

where  $\beta = (\beta_1, \dots, \beta_m)$  is a vector of “treatment effect”s, where  $\beta_j$  is the effect at gene  $j$ , and  $\sigma = (\sigma_1, \dots, \sigma_m)$  is a vector of standard deviation parameters. For each gene  $j$  we wish to test the null  $H_j : \beta_j = 0$  (that is, that there is no treatment effect). For simplicity you can assume that  $\sigma_j = 1$  is known for all  $j$ . You can also assume that  $n = 10$  and  $m = 1,000$ .

Assume that the true effects  $\beta_j$  are independent, and identically distributed, with

$$\beta_j \sim \pi_0 \delta_0 + (1 - \pi_0) N(0, \sigma_b^2). \quad (2)$$

where  $\delta_0$  denotes a point mass on zero. That is,  $\beta_j = 0$  with probability  $\pi_0$ , and  $\beta_j \sim N(0, \sigma_b^2)$  with probability  $1 - \pi_0$ .

- i) Explain how (1) and (2) together define a hierarchical model for the data and treatment effects at each gene. Given  $\pi_0, \sigma_b$ , are the  $\beta_j$  exchangeable? What about unconditional on  $\pi_0, \sigma_b$ ? What about the data matrix  $D_{ij}$  - which aspects of this matrix are exchangeable, if any? Explain your reasoning.
- ii) Write an R function to simulate data  $D$  under this hierarchical model, for user-specified  $\pi_0$  and  $\sigma_b$ . The function should take  $\pi_0$  and  $\sigma_b$  as input, and return a list, with elements  $D$  (a matrix) and  $\beta$  (a vector).

- iii) Write an R function to compute a  $p$  value  $p_j$  for each column of the data matrix  $D$ , testing  $H_0 : \beta_j = 0$ . This function should take as input the data matrix  $D$  and output a vector of  $p$  values. You can use any reasonable two-sided test, but state which test you use. Apply your R function to data simulated under a)  $\pi_0 = 1$ , b)  $\pi_0 = 0.5, \sigma_b = 3$ ; c)  $\pi_0 = 0, \sigma_b = 3$ . Provide histograms of the  $p$  values in each case and comment on their distributions.
  - iv) Write an R function to apply the Benjamini–Hochberg rule to control FDR at a user-specified level  $\alpha$ . This function should input a vector of  $p$  values, and a level  $\alpha$ , and output a vector of binary (0/1) indicators,  $\gamma = (\gamma_1, \dots, \gamma_m)$  say, where  $\gamma_j = 1$  indicates that the rule would reject  $H_j : \beta_j = 0$ .
  - v) Write an R function to compute the empirical False Discovery Rate (i.e. the number of false discoveries divided by the number of discoveries) for any given value for the vector  $\beta$  of true values of  $\beta$ , and the vector  $\gamma$  of reject decisions. That is, the function should return  $V/R$  in the notation of the notes. Remember to deal correctly with the special case of no discoveries,  $R = 0$ .
  - vi) Perform a simulation study to estimate the actual FDR ( $E(V/R)$ ) achieved by the BH rule in the three cases a), b) and c) above. In each case perform the test procedure for different levels  $\alpha$ , and plot the estimated  $E(V/R)$  as a function of  $\alpha$  (say for  $\alpha = (0.05, 0.1, \dots, 0.5)$ ). Comment on the results. [NOTE: to estimate the actual FDR you have to estimate  $E(V/R)$  where the expectation is over datasets  $D$ . To do this you will want to do a simulation study where you simulate a large number of datasets  $D$ , not just one dataset!]
  - vii) Repeat the simulation study, but this time estimate the pFDR instead of the FDR, and plot this as a function of  $\alpha$ .
2. The `qvalue` package in R implements Storey's approach to estimating FDR. To install this package use

```
source("http://bioconductor.org/biocLite.R")
biocLite("qvalue")
library("qvalue")
```

The package takes a vector of  $p$  values, and outputs a list which includes an estimate of  $\pi_0$  (obtained using the  $p$  values near 1) and a vector of  $q$  values. Try, for example, for a vector of  $p$  values  $\mathbf{p}$ ,

```
res=qvalue(p)
res$pi0
res$qvalues
```

The  $q$  value for a particular observation is an estimate of the pFDR if you reject all things that are as or more significant than that observation. You can convert the vector of  $q$  values into a list of reject decisions at a given  $\alpha$  level (the  $\gamma$  vector above) using, say,

```
compute.gamma=function(q,alpha){return(q<alpha)}
```

- i) Repeat the simulation study above, using `qvalue` instead of the BH procedure. Produce plots of the FDR vs the  $\alpha$  level for `qvalue` and compare them with those obtained for BH.
  - ii) Perform a simulation study (e.g. by modifying the simulations you have already performed), to see how accurately `qvalue` is able to estimate the proportion of nulls  $\pi_0$ . Try varying  $\pi_0$  from 0 to 1 for at least 3 different values of  $\sigma_b$ , and in each case provide plots of the true  $\pi_0$  vs the estimated  $\pi_0$  from `qvalue`. Comment on the results.
3. Now consider implementing an Empirical Bayes approach to this problem. To do so, given data  $D$  we will need two steps:
- A Estimate the hyper parameters  $\pi_0, \sigma_b$  in (2) by maximum likelihood. Call the estimates  $\hat{\pi}_0, \hat{\sigma}_b$ .
  - B Compute the posterior distribution  $p(\beta_j | D, \hat{\pi}_0, \hat{\sigma}_b)$  for each  $j$ .

This question takes you through these two steps.

- i) Define  $\bar{D}_j = (1/n) \sum_i D_{ij}$ . Show that the vector  $\bar{D} := (\bar{D}_1, \dots, \bar{D}_m)$  is sufficient for  $\beta$ . That is,  $p(D|\beta) \propto p(\bar{D}|\beta)$  where the constant of proportionality does not depend on  $\beta$ . [This means that, as far as inference for  $\beta$  is concerned, the likelihood  $p(D|\beta)$  is equivalent to the likelihood  $p(\bar{D}|\beta)$ , so from now on you can treat  $\bar{D}$  as your data instead of  $D$ .]

- ii) Derive an expression for the log-likelihood  $l(\pi_0, \sigma_b) := \log(p(\bar{D}|\pi_0, \sigma_b))$ .  
[Hint: note that the  $\bar{D}_j$  are independent given  $\pi_0, \sigma_b$ ]
- iii) Write an R function to compute the log-likelihood  $l(\pi_0, \sigma_b)$ , or alternatively  $l(\theta_1, \theta_2)$  where  $\theta_1 = \log(\pi_0/(1 - \pi_0))$ ,  $\theta_2 = \log(\sigma_b)$ .  
[The motivation for this reparameterization is that  $\theta_1, \theta_2$  can take any value on the real line.] Try using the R function optimize (or another method if you prefer) to maximize the likelihood over  $\pi_0, \sigma_b$  (or  $\theta_1, \theta_2$ ). [You may or may not find that this works... it is a somewhat tricky numerical problem. The reparameterization may help. Alternatively if you know about the EM algorithm you can try that.]
- iv) Derive the posterior distribution  $\beta_j|D, \pi_0, \sigma_b$ . Hint: this posterior should be a mixture of a point mass at zero and a normal distribution. It may help to first derive  $p(\beta_j = 0|D, \pi_0, \sigma_b)$ , and then  $p(\beta_j|D, \pi_0, \sigma_b, \beta_j \neq 0)$ .
- v) Implement a method that computes  $p(\beta_j = 0|D, \hat{\pi}_0, \hat{\sigma}_b)$ . Implement another method that takes these probabilities and rejects those tests  $j$  for which this probability is  $< \alpha$ . Add this method to your simulation study and see how it performs. (If you are unable to get the optimization for  $\pi, \sigma_b$  to work then you can “cheat” in this step and use the true value of  $\pi, \sigma_b$ .)